

## Longest increasing subsequence

We are given an array with  $n$  numbers:  $a[0 \dots n - 1]$ . The task is to find the longest, strictly increasing, subsequence in  $a$ .

Formally we look for the longest sequence of indices  $i_1, \dots i_k$  such that

$$i_1 < i_2 < \dots < i_k, \quad a[i_1] < a[i_2] < \dots < a[i_k]$$

In this article we discuss multiple algorithms for solving this task. Also we will discuss some other problems, that can be reduced to this problem.

### Solution in $O(n^2)$ with dynamic programming

Dynamic programming is a very general technique that allows to solve a huge class of problems. Here we apply the technique for our specific task.

First we will search only for the **length** of the longest increasing subsequence, and only later learn how to restore the subsequence itself.

#### Finding the length

To accomplish this task, we define an array  $d[0 \dots n - 1]$ , where  $d[i]$  is the length of the longest increasing subsequence that ends in the element at index  $i$ .

##### Example

$$\begin{aligned} a &= \{8, 3, 4, 6, 5, 2, 0, 7, 9, 1\} \\ d &= \{1, 1, 2, 3, 3, 1, 1, 4, 5, 2\} \end{aligned}$$

The longest increasing subsequence that ends at index 4 is  $\{3, 4, 5\}$  with a length of 3, the longest ending at index 8 is either  $\{3, 4, 5, 7, 9\}$  or  $\{3, 4, 6, 7, 9\}$ , both having length 5, and the longest ending at index 9 is  $\{0, 1\}$  having length 2.

We will compute this array gradually: first  $d[0]$ , then  $d[1]$ , and so on. After this array is computed, the answer to the problem will be the maximum value in the array  $d[]$ .

So far we only learned how to find the length of the subsequence, but not how to find the subsequence itself.

To be able to restore the subsequence we generate an additional auxiliary array  $p[0 \dots n - 1]$  that we will compute alongside the array  $d[]$ .  $p[i]$  will be the index  $j$  of the second last element in the longest increasing subsequence ending in  $i$ . In other words the index  $p[i]$  is the same index  $j$  at which the highest value  $d[j]$  was obtained. This auxiliary array  $p[]$  points in some sense to the ancestors.

Then to derive the subsequence, we just start at the index  $i$  with the maximal  $d[i]$ , and follow the ancestors until we deduced the entire subsequence, i.e. until we reach the element with  $d[i] = 1$ .

#### Implementation of restoring

We will change the code from the previous sections a little bit. We will compute the array  $p[]$  alongside  $d[]$ , and afterwards compute the subsequence.

For convenience we originally assign the ancestors with  $p[i] = -1$ . For elements with  $d[i] = 1$ , the ancestors value will remain  $-1$ , which will be slightly more convenient for restoring the subsequence.

```
vector<int> lis(vector<int> const& a) {
    int n = a.size();
    vector<int> d(n, 1), p(n, -1);
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < i; j++) {
            if (a[j] < a[i] && d[j] < d[i] + 1) {
                d[i] = d[j] + 1;
                p[i] = j;
            }
        }
    }

    int ans = d[0], pos = 0;
    for (int i = 1; i < n; i++) {
        if (d[i] > ans) {
            ans = d[i];
            pos = i;
        }
    }

    vector<int> subseq;
    while (pos != -1) {
        subseq.push_back(a[pos]);
        pos = p[pos];
    }
    reverse(subseq.begin(), subseq.end());
    return subseq;
}
```

#### Alternative way of restoring the subsequence

It is also possible to restore the subsequence without the auxiliary array  $p[]$ . We can simply recalculate the current value of  $d[i]$  and also see how the maximum was reached.

So let the current index be  $i$ . I.e. we want to compute the value  $d[i]$  and all previous values  $d[0], \dots, d[i - 1]$  are already known. Then there are two options:

- $d[i] = 1$ : the required subsequence consists only of the element  $a[i]$ .
- $d[i] > 1$ : the subsequence will end at  $a[i]$ , and right before it will be some number  $a[j]$  with  $j < i$  and  $a[j] < a[i]$ .

It's easy to see, that the subsequence ending in  $a[j]$  will itself be one of the longest increasing subsequences that ends in  $a[j]$ . The number  $a[i]$  just extends that longest increasing subsequence by one number.

Therefore, we can just iterate over all  $j < i$  with  $a[j] < a[i]$ , and take the longest sequence that we get by appending  $a[i]$  to the longest increasing subsequence ending in  $a[j]$ . The longest increasing subsequence ending in  $a[j]$  has length  $d[j]$ , extending it by one gives the length  $d[j] + 1$ .

$$d[i] = \max_{\substack{j < i \\ a[j] < a[i]}} (d[j] + 1)$$

If we combine these two cases we get the final answer for  $d[i]$ :

$$d[i] = \max \left( 1, \max_{\substack{j < i \\ a[j] < a[i]}} (d[j] + 1) \right)$$

#### Implementation

Here is an implementation of the algorithm described above, which computes the length of the longest increasing subsequence.

```
int lis(vector<int> const& a) {
    int n = a.size();
    vector<int> d(n, 1);
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < i; j++) {
            if (a[j] < a[i])
                d[i] = max(d[i], d[j] + 1);
        }
    }

    int ans = d[0];
    for (int i = 1; i < n; i++) {
        ans = max(ans, d[i]);
    }
    return ans;
}
```

#### Restoring the subsequence

This method leads to a slightly longer code, but in return we save some memory.

### Solution in $O(n \log n)$ with dynamic programming and binary search

In order to obtain a faster solution for the problem, we construct a different dynamic programming solution that runs in  $O(n^2)$ , and then later improve it to  $O(n \log n)$ .

We will use the dynamic programming array  $d[0 \dots n]$ . This time  $d[l]$  doesn't correspond to the element  $a[i]$  or to an prefix of the array.  $d[l]$  will be the smallest element at which an increasing subsequence of length  $l$  ends.

Initially we assume  $d[0] = -\infty$  and for all other lengths  $d[l] = \infty$ .

We will again gradually process the numbers, first  $a[0]$ , then  $a[1]$ , etc, and in each step maintain the array  $d[]$  so that it is up to date.

##### Example

Given the array  $a = \{8, 3, 4, 6, 5, 2, 0, 7, 9, 1\}$ , here are all their prefixes and their dynamic programming array. Notice, that the values of the array don't always change at the end.

prefix = {}	$d = \{-\infty, \infty, \dots\}$
prefix = {8}	$d = \{-\infty, 8, \infty, \dots\}$
prefix = {8, 3}	$d = \{-\infty, 3, \infty, \dots\}$
prefix = {8, 3, 4}	$d = \{-\infty, 3, 4, \infty, \dots\}$
prefix = {8, 3, 4, 6}	$d = \{-\infty, 3, 4, 6, \infty, \dots\}$
prefix = {8, 3, 4, 6, 5}	$d = \{-\infty, 3, 4, 5, \infty, \dots\}$
prefix = {8, 3, 4, 6, 5, 2}	$d = \{-\infty, 2, 4, 5, \infty, \dots\}$
prefix = {8, 3, 4, 6, 5, 2, 0}	$d = \{-\infty, 0, 4, 5, \infty, \dots\}$
prefix = {8, 3, 4, 6, 5, 2, 0, 7}	$d = \{-\infty, 0, 4, 5, 7, \infty, \dots\}$
prefix = {8, 3, 4, 6, 5, 2, 0, 7, 9}	$d = \{-\infty, 0, 4, 5, 7, 9, \infty, \dots\}$
prefix = {8, 3, 4, 6, 5, 2, 0, 7, 9, 1}	$d = \{-\infty, 0, 1, 5, 7, 9, \infty, \dots\}$

When we process  $a[i]$ , we can ask ourselves. What have the conditions to be, that we write the current number  $a[i]$  into the  $d[0 \dots n]$  array?

We set  $d[l] = a[i]$ , if there is a longest increasing sequence of length  $l$  that ends in  $a[i]$ , and there is no longest increasing sequence of length  $l$  that ends in a smaller number. Similar to the previous approach, if we remove the number  $a[i]$  from the longest increasing sequence of length  $l$ , we get another longest increasing sequence of length  $l - 1$ . So we want to extend a longest increasing sequence of length  $l - 1$  by the number  $a[i]$ , and obviously the longest increasing sequence of length  $l - 1$  that ends with the smallest element will work the best, in other words the sequence of length  $l - 1$  that ends in element  $d[l - 1]$ .

There is a longest increasing sequence of length  $l - 1$  that we can extend with the number  $a[i]$ , exactly if  $d[l - 1] < a[i]$ . So we can just iterate over each length  $l$ , and check if we can extend a longest increasing sequence of length  $l - 1$  by checking the criteria.

Additionally we also need to check, if we maybe have already found a longest increasing sequence of length  $l$  with a smaller number at the end. So we only update if  $a[i] < d[l]$ .

After processing all the elements of  $a[]$  the length of the desired subsequence is the largest  $l$  with  $d[l] < \infty$ .

```
int lis(vector<int> const& a) {
    int n = a.size();
    const int INF = 1e9;
    vector<int> d(n+1, INF);
    d[0] = -INF;

    for (int i = 0; i < n; i++) {
        for (int l = 1; l <= n; l++) {
            if (d[l-1] < a[i] && a[i] < d[l])
                d[l] = a[i];
        }

        int ans = 0;
        for (int l = 0; l <= n; l++) {
            if (d[l] < INF)
                ans = l;
        }
        return ans;
    }
}
```

We now make two important observations.

- 1. The array  $d$  will always be sorted:  $d[l - 1] < d[l]$  for all  $i = 1 \dots n$ .  
This is trivial, as you can just remove the last element from the increasing subsequence of length  $l$ , and you get a increasing subsequence of length  $l - 1$  with a smaller ending number.
- 2. The element  $a[i]$  will only update at most one value  $d[l]$ .  
This follows immediately from the above implementation. There can only be one place in the array with  $d[l - 1] < a[i] < d[l]$ .

Thus we can find this element in the array  $d[]$  using [binary search](#) in  $O(\log n)$ . In fact we can simply look in the array  $d[]$  for the first number that is strictly greater than  $a[i]$ , and we try to update this element in the same way as the above implementation.

Implementation

This gives us the improved  $O(n \log n)$  implementation:

```
int lis(vector<int> const& a) {
    int n = a.size();
    const int INF = 1e9;
```

from  $0$  to  $n - 1$ ), or use a dynamic segment tree (only generate the branches of the tree that are important). Otherwise the memory consumption will be too high.

On the other hand this method has also some **advantages**: with this method you don't have to think about any tricky properties in the dynamic programming solution. And this approach allows us to generalize the problem very easily (see below).

Related tasks

Here are several problems that are closely related to the problem of finding the longest increasing subsequence.

Longest non-decreasing subsequence

This is in fact nearly the same problem. Only now it is allowed to use identical numbers in the subsequence.

The solution is essentially also nearly the same. We just have to change the inequality signs, and make a slightly modification to the binary search.

Number of longest increasing subsequences

We can use the first discussed method, either the  $O(n^2)$  version or the version using data structures. We only have to additionally store in how many ways we can obtain longest increasing subsequences ending in the values  $d[i]$ .

The number of ways to form a longest increasing subsequences ending in  $a[i]$  is the sum of all ways for all longest increasing subsequences ending in  $j$  where  $d[j]$  is maximal. There can be multiple such  $j$ , so we need to sum all of them.

Using a Segment tree this approach can also be implemented in  $O(n \log n)$ .

It is not possible to use the binary search approach for this task.

Smallest number of non-increasing subsequences covering a sequence

For a given array with  $n$  numbers  $a[0 \dots n - 1]$  we have to colorize the numbers in the smallest number of colors, so that each color forms a non-increasing subsequence.

To solve this, we notice that the minimum number of required colors is equal to the length of the longest increasing subsequence.

**Proof:** We need to prove the **duality** of these two problems.

Let's denote by  $x$  the length of the longest increasing subsequence and by  $y$  the least number of non-increasing subsequences that form a cover. We need to prove that  $x = y$ .

```
vector<int> d(n+1, INF);
d[0] = -INF;

for (int i = 0; i < n; i++) {
    int l = upper_bound(d.begin(), d.end(), a[i]) - d.begin();
    if (d[l-1] < a[i] && a[i] < d[l])
        d[l] = a[i];
}

int ans = 0;
for (int l = 0; l <= n; l++) {
    if (d[l] < INF)
        ans = l;
}
return ans;
}
```

Restoring the subsequence

It is also possible to restore the subsequence using this approach. This time we have to maintain two auxiliary arrays. One that tells us the index of the elements in  $d[]$ . And again we have to create an array of "ancestors"  $p[i]$ .  $p[i]$  will be the index of the previous element for the optimal subsequence ending in element  $i$ .

It's easy to maintain these two arrays in the course of iteration over the array  $a[]$  alongside the computations of  $d[]$ . And at the end it is not difficult to restore the desired subsequence using these arrays.

Solution in  $O(n \log n)$  with data structures

Instead of the above method for computing the longest increasing subsequence in  $O(n \log n)$  we can also solve the problem in a different way: using some simple data structures.

Let's go back to the first method. Remember that  $d[i]$  is the value  $d[j] + 1$  with  $j < i$  and  $a[j] < a[i]$ .

Thus if we define an additional array  $t[]$  such that

$$t[a[i]] = d[i],$$

then the problem of computing the value  $d[i]$  is equivalent to finding the **maximum value in a prefix** of the array  $t[]$ :

$$d[i] = \max(t[0 \dots a[i] - 1] + 1)$$

The problem of finding the maximum of a prefix of an array (which changes) is a standard problem that can be solved by many different data structures. For instance we can use a [Segment tree](#) or a [Fenwick tree](#).

This method has obviously some **shortcomings**: in terms of length and complexity of the implementation this approach will be worse than the method using binary search. In addition if the input numbers  $a[i]$  are especially large, then we would have to use some tricks, like compressing the numbers (i.e. renumber them